



Getting Started on Kestrel

April 2024

Haley Yandt & Olivia Hull

Introduction

Background

- NREL provides HPC in support of the DOE Office of Energy Efficiency and Renewable Energy (EERE)
 - NREL runs yearly open call (June) on behalf of EERE
 - NREL HPC is the principal provider of high performance and data-intensive computing services to the DOE EERE Offices
 - additional call this year (student applications), allocating 5.6M AU (560k node hours)
- **EERE Mission:** ...accelerate the research, development, demonstration, and deployment of technologies and solutions to equitably transition America to net-zero greenhouse gas emissions...(<https://www.energy.gov/eere/mission>)
- **NREL Mission:** NREL strives to achieve our vision of a clean energy future for the world through our mission: leading research, innovation, and strategic partnerships to deliver solutions for a clean energy economy...(<https://www.nrel.gov/about/vision.html>)

Introduction

Allocations & Accounts

Allocation Questions: HPC-Requests@nrel.gov

General HPC Questions: HPC-Help@nrel.gov

- ‘some’ allocations activated Apr 22
 - defaulted to 100k AU (10k node hours) for CPU requests
 - contact **HPC-Requests** if additional AUs needed (up to 400k max)
 - remainder to be activated when account process complete
- accounts activated or in progress
 - additional paperwork required in some cases, please respond promptly to requests
 - reach out to **HPC-Help** if you need assistance with your account
- software licenses
 - unable to provide licenses in most cases
 - follow documentation for requesting access on webpage for listed software
 - reach out to HPC-Help for all other licensing issues
 - software docs: <https://nrel.github.io/HPC/Documentation/Applications/>

Connect to Kestrel

- Connect directly to Kestrel login node

```
$ ssh <username>@kestrel.nrel.gov # Login externally
...
Password+OTPToken:
```

- Alternatively, you can use the [HPC VPN](#) or the [SSH gateway](#).
- Logins nodes serve as a gateway to the rest of the system and are shared resources.

Data Analysis and Visualization

- Kestrel has 8 DAV nodes, intended for HPC applications that require a graphical user interface.
 - One DAV node accessible directly to non-NREL users at **kestrel-dav.nrel.gov**
 - The other 7 are accessible via the HPC VPN
- [FastX](#) is available for a graphical remote desktop
- Can connect to DAV nodes using FastX through a web or desktop client
- DAV nodes are not general-purpose remote desktops. They are intended for HPC or visualization software that requires Kestrel.
- DAV nodes are shared resources. Computationally intensive work needs to be done on compute nodes.

System Configuration

- **44 Pflops peak**
- Interconnect: HPE Slingshot
- Compute nodes:
 - 2144 Dual socket Intel Sapphire Rapids 52-core processors (**104 cores** total per node)
 - Standard nodes have **256GB RAM**
 - 256 nodes have **1.75 NVMe** local disk
 - **10 bigmem nodes** with **2TB RAM** and **5.5TB NVMe** local disk
- GPU nodes:
 - Coming soon – in final stages of our acceptance testing
 - 132 GPU nodes with 4 NVIDIA H100 GPUs per node
- Additional information: <https://www.nrel.gov/hpc/kestrel-system-configuration.html>

Kestrel Filesystem

- **95 PB** ClusterStor Lustre Filesystem
- Parallel File System (PFS)
 - Intended for high performance I/O
 - /projects/<allocation handle>
 - /scratch/\$USER
 - Run jobs out of /scratch, use /projects for longer term storage and sharing data with allocation team members
 - No quota on /scratch, but subject to purging after 28 days of inactivity
 - There are no backups of PFS data
- Home File System: 1.2PB, part of the ClusterStor system
 - User /home directories have a 50GB quota
 - Shouldn't run jobs out of /home - not tuned for parallel I/O
- [Globus](#) is available for large data transfers



Running Jobs

- Kestrel uses the [Slurm](#) job scheduler
- Submitting jobs:
 - **sbatch**: Use to submit a batch job
 - **salloc**: Used to allocate resources for an interactive job
- Slurm includes a [suite of command-line tools](#) to monitor jobs in the queue
- Job priority is calculated as a sum of multiple factors including age, jobsize, QOS, and fairshare
- Quality of service (QOS) :
 - standby: jobs with standby qos will only run if there are idle nodes that non-standby jobs can't use. No AU cost. Default for allocations with no AUs remaining, or can also specify in your job script.
 - **--qos=high**: gives small priority bump. Allocation will be charged at twice the standard rate of AUs
- Fairshare: calculated based on quarter AU assignment/total system AUs
 - Those running more than their fair share in the past two weeks will have decreased priority
 - Large impact on priority calculation
- Required flags for your jobs
 - **--account=<project handle>**: The allocation to use for the job
 - **--time=<walltime>**: Maximum Job Duration (walltime)

Partitions

- Compute nodes belong to one or more partitions. A partition is a collection of compute nodes which often share similar characteristics (memory, gpu, etc.).
- [Kestrel partition specifications](#)
- In general, we recommend not specifying a partition. Slurm automatically routes jobs to the appropriate partition based on requested resources.
 - Exceptions are debug and shared partitions
- `--partition=debug`: Nodes dedicated to developing and troubleshooting jobs
- `--partition=shared`
 - Nodes in the shared partition can be shared by multiple users or jobs. This partition is intended for jobs that do not require a whole node.
 - These nodes have 250GB of usable RAM and 104 cores. By default, your job will be allocated 1.024GB of RAM per core requested. You can use `--mem` or `--mem-per-cpu` to change this.
 - AU cost is calculated based on either the amount of cores or the amount of memory requested, whichever is a greater percentage of the total of that resource available on the node.

Accounting

- Track AU usage at <https://hpcprojects.nrel.gov>
 - Login with HPC account credentials
- AU cost = (Walltime in hours * Number of Nodes * QoS Factor * Charge Factor)
- Charge factor on Kestrel is 10 for CPU nodes, and 25 per GPU (100 per GPU node)

Allocation Unit (AU) Use Report for csc000

[Summary of Request](#)[View Award Letter](#)[AU Use Report](#)[List Jobs](#)

Project Use

System	AUs Requested	AUs Assigned	Normal AUs Charged	High-Pri AUs Charged	Buy-in AUs Charged	Standby AUs Used	Total AUs Charged	Total AUs Used	AUs Planned To Date	AUs Remaining
Eagle	—	50,000	24,450	0	0	42	24,450	24,493	27,198	25,550
Swift	—	50,000	9,713	0	0	0	9,713	9,713	27,198	40,287
Vermilion	—	50,000	39	0	0	0	39	39	27,198	49,961
Kestrel	—	20,000	301,014	0	0	127	301,014	301,141	10,879	0
ALL	20,000	170,000	335,216	0	0	169	335,216	335,385	92,473	115,798

Example submit script

```
#!/bin/bash
#SBATCH --nodes=2
#SBATCH --tasks-per-node=104
#SBATCH --time=2:00:00
#SBATCH --account=<your-account-name>
#SBATCH --job-name=<your-job-name>
```

```
module load vasp/<version>
```

```
srun vasp_std |& tee out
```

Non-shared partition

```
#!/bin/bash
#SBATCH --nodes=1
#SBATCH --tasks-per-node=26
#SBATCH --time=2:00:00
#SBATCH --partition=shared
#SBATCH --mem-per-cpu=2G
#SBATCH --account=<your-account-name>
#SBATCH --job-name=<your-job-name>
```


```
module load vasp/<version>
```

```
srun vasp_std |& tee out
```

Shared partition

Note the --mem-per-cpu directive

Environments



- Kestrel multiprocessor: Intel Sapphire Rapids, 52 CPU cores
- Kestrel node: two multiprocessors, **104 CPU cores total per node**
- Kestrel Cluster: 2100+ CPU nodes
- Hardware informs compiler/MPI choices

The "optimal environment" depends on your needs, but for CPU nodes:

- Intel processors → Intel compilers
- Cray Slingshot network → Cray MPICH or Cray MPICH ABI (multinode)

Using Anaconda

- Load the anaconda module:
 `module load anaconda3`
- Mamba is also available (install packages much faster)
- Install conda environments somewhere other than home directory
 - For example, project directory: `/projects/your_project_name`
- Mpi4py in anaconda utilizes OpenMPI, which is unstable for multinode jobs on our machine. If you need to use mpi4py multinode, please reach out to hpc-help@nrel.gov
- Documentation
link: <https://nrel.github.io/HPC/Documentation/Environment/Customization/conda/>

Pre-made modules

- Use the command ``module avail [name]`` to see available modules
- To see more information on a module, use the command ``module show [name]``.
- For example, if you are using a code like vasp, all you will need in order to set up your environment is ``module load vasp/<version>`` (and you can see which versions are available with ``module avail vasp``)
- Some applications have NREL documentation pages with sample submit scripts: <https://nrel.github.io/HPC/Documentation/Applications/>
 - Not an exhaustive list

PrgEnvs

- If you have not worked on a Cray machine before, environments and modules may work differently than you expect.
- Default environment is 'PrgEnv-cray'
- PrgEnvs are "modules of modules" - when you load a PrgEnv, you load a bundle of modules. For PrgEnv-cray:
 - Cray compiler environment "CCE" compilers (C, C++, Fortran)
 - Cray-mpich: Cray's MPI package, works better than other MPIs on our network
 - Some underlying communication/compilation libraries
 - Can see list of modules associated with PrgEnv-cray with ``module show PrgEnv-cray``
- We recommend PrgEnv-cray (or PrgEnv-intel) for CPU codes.
- Tutorial on PrgEnvs:
<https://nrel.github.io/HPC/Documentation/Systems/Kestrel/Environments/tutorial/>

Compiling code within PrgEnvs

- The PrgEnvs work differently than "typical" module environments
 - use the Cray wrappers `cc`, `CC`, and `ftn` (for `c`, `c++`, `fortran`, respectively)
- Wrapper on top of both the compiler and the MPI
- But flags you pass should be for underlying compiler
- May need to load additional modules like `cray-fftw`, `cray-hdf5-parallel`, etc.

Env	Wrappers	Underlying compiler	MPI	MPI location
PrgEnv-cray	<code>cc</code> , <code>CC</code> , <code>ftn</code>	CCE suite	Cray MPICH	<code>/opt/cray/pe/mpich/8.1.23/ofi/crayclang/10.0/bin/</code>
PrgEnv-intel	<code>cc</code> , <code>CC</code> , <code>ftn</code>	Intel suite (<code>ifort</code> , etc)	Cray MPICH	<code>/opt/cray/pe/mpich/8.1.23/ofi/intel/19.0/bin/</code>
PrgEnv-gnu	<code>cc</code> , <code>CC</code> , <code>ftn</code>	GNU (<code>gcc</code> , <code>gfortran</code> , etc)	Cray MPICH	<code>/opt/cray/pe/mpich/8.1.23/ofi/gnu/9.1/bin/</code>

```
cc --version
```

```
cc -o hello_world_mpi.exe hello_world_mpi.c
```


Environments

- The default environments are PrgEnvs
 - PrgEnv-cray, PrgEnv-intel, PrgEnv-gnu...
 - Compilers given in name
 - e.g. PrgEnv-intel uses intel compilers
 - MPI used is **Cray MPICH**
- Why use PrgEnv-*?
 - Cray MPICH better utilizes the Cray Slingshot network, so multi-node jobs run faster than with other MPIs
- NREL-built environments are also available
 - These behave similar to those on Eagle
- Refer to the environments tutorial for instructions on swapping between "realms" <https://nrel.github.io/HPC/Documentation/Systems/Kestrel/Environments/tutorial/>

PrgEnv "Realm"

<u>Compilers:</u>	<u>MPI:</u>	<u>Scientific Libraries:</u>
Cray CCE	Cray MPICH	cray-libsci
Intel		cray-fftw
GCC		cray-hdf5-parallel

/opt/cray/

Example:

module load PrgEnv-gnu
- Loads gcc compilers, Cray MPICH, cray-libsci, a bundle of communication libraries

NREL-built "Realm"

<u>Compilers:</u>	<u>MPI:</u>	<u>Scientific Libraries:</u>
Intel	Intel MPI	MKL, GSL
GCC	Open MPI MPICH	Open BLAS, LaPACK, HDF5, etc.

/nopt/nrel/apps/modules/

Example:

module load hdf5/1.14.1-2-
openmpi-gcc
- HDF5 built with gcc
compilers and OpenMPI

Compiling

- Want to compile a code that needs MPI, scalapack, fftw, hdf5

PrgEnv-cray Example:

```
module restore  
module load cray-libsci  
module load cray-fftw  
module load cray-hdf5-parallel
```

What this gets you:

Program built with Cray compilers, Cray MPICH, and Cray's implementations of fftw (auto-loads correct build), scalapack (libsci), and hdf5

NREL Envs Example (intel toolchain):

```
module restore  
module unload PrgEnv-cray  
module unload cce  
module load intel-oneapi-compilers  
module load intel-oneapi-mpi  
module load intel-oneapi-mkl  
module load fftw/3.3.10-intel-oneapi-mpi-intel  
module load hdf5/1.14.1-2-intel-oneapi-mpi-intel
```

What this gets you:

Program built with intel compilers, Intel MPI, scalapack in MKL, fftw and hdf5 built with intel/intel MPI

Quick Reference Page

- To connect:

```
[user@mac ~]$ ssh <username>@kestrel.hpc.nrel.gov
```

- PrgEnv-cray loaded by default
- To set up for NREL-built environment modules:

```
[user@kestrel ~]$ module restore  
[user@kestrel ~]$ module unload PrgEnv-cray
```

- If using a PrgEnv:
 - use the Cray wrappers (ftn/cc/CC) to compile
 - module load cray-[fftw, libsci, hdf5-parallel, etc] for scientific libraries
- “salloc -N 1 -n 104 --time=1:00:00 --account=[your account]” for an interactive session
- Build codes on compute nodes (via salloc, etc.)
- For applications that were compiled with OpenMP threading enabled, include in slurm script:
export OMP_PROC_BIND=spread or export KMP_AFFINITY=balanced (intel)
- Getting started on Kestrel docs: https://nrel.github.io/HPC/Documentation/Systems/Kestrel/getting_started_kestrel/

Issues

- **MPI Scaling**
 - Poor scaling in multinode jobs is a known issue
 - We are working to get this resolved as fast as possible
 - See performance recommendations here: [https://nrel.github.io/HPC/Documentation/Systems/eagle to kestrel transition/#5-performance-recommendations](https://nrel.github.io/HPC/Documentation/Systems/eagle%20to%20kestrel%20transition/#5-performance-recommendations)
- We don't expect OpenMPI to perform multinode
- Reporting issues
 - Submit tickets to hpc-help@nrel.gov
 - Include details on the problem
 - Include job ID and any relevant outputs and inputs (log files, std-err, std-out, etc.)
 - Modules loaded, sbatch file, etc.

Resources

- Kestrel environments resources:
 - Overview of environments on Kestrel:
<https://nrel.github.io/HPC/Documentation/Systems/Kestrel/Environments/>
 - Very simple
tutorial: https://github.com/NREL/HPC/tree/master/kestrel/mpi_version_check
 - More complex environments
tutorial: <https://nrel.github.io/HPC/Documentation/Systems/Kestrel/Environments/tutorial/>
- Code (examples, etc.) repo: <https://github.com/NREL/HPC/tree/master/kestrel>
- HPC office hours
 - <https://www.nrel.gov/hpc/training.html>
 - Stop by with questions
- Tickets: hpc-help@nrel.gov
- Documentation site: <https://nrel.github.io/HPC/Documentation/>

Feedback is Appreciated!

If you have any suggestions to improve this presentation we invite you to share with us at HPC-Help@nrel.gov

Thank You

www.nrel.gov

NREL is a national laboratory of the U.S. Department of Energy, Office of Energy Efficiency and Renewable Energy, operated by the Alliance for Sustainable Energy, LLC.

